

Heart Disease Prediction -- End-to-End MLOps Report

Author: Harsh

Repository: <https://github.com/HARSH6355/Heart-Disease-Predictions-MLOPS>

Date: July 2026

Table of Contents

1. Introduction & Objective
 2. Dataset Description
 3. System Architecture
 4. Environment Setup & Installation
 5. Exploratory Data Analysis (EDA)
 6. Feature Engineering & Preprocessing
 7. Modelling Choices & Evaluation
 8. Experiment Tracking with MLflow
 9. Flask Prediction API & Web Interface
 10. Unit Testing
 11. CI/CD Pipeline -- GitHub Actions
 12. Docker Containerization
 13. Kubernetes Deployment
 14. Monitoring with Prometheus
 15. Conclusion & Future Work
-

1. Introduction & Objective

Heart disease remains one of the leading causes of death globally, accounting for approximately 17.9 million deaths each year according to the World Health Organization. Early detection and risk assessment are critical for timely intervention and improved patient outcomes.

The objective of this project is to design, develop, and deploy a scalable, reproducible, and monitored end-to-end Machine Learning classification solution to predict the risk of heart disease based on patient health data. The workflow adheres to modern MLOps best practices encompassing:

- Automated data ingestion from a trusted public repository
- Robust feature transformation using scikit-learn pipelines

- Experiment tracking with MLflow for full reproducibility
- Model serving via a Flask REST API with a polished web interface
- Unit testing with Pytest
- Continuous Integration via GitHub Actions
- Containerized deployment using Docker
- Orchestration with Kubernetes manifests
- Production monitoring with Prometheus metrics

This report documents every stage of the pipeline, from raw data to production deployment, with supporting screenshots and technical detail.

2. Dataset Description

Property	Value
Name	Heart Disease UCI Dataset
Source	UCI Machine Learning Repositor
Instances	303
Features	13 (categorical, integer, and
Target	Binary (0 = No heart disease,

2.1 Feature Dictionary

Feature	Description	Type
`age`	Age in years	Integer
`sex`	Sex (1 = Male, 0 = Female)	Categorical
`cp`	Chest pain type (1-4)	Categorical
`trestbps`	Resting blood pressure (mm Hg)	Integer
`chol`	Serum cholesterol (mg/dl)	Integer
`fbs`	Fasting blood sugar > 120 mg/d	Categorical
`restecg`	Resting ECG results (0-2)	Categorical
`thalach`	Maximum heart rate achieved	Integer
`exang`	Exercise-induced angina (1 = Y	Categorical
`oldpeak`	ST depression induced by exerc	Float
`slope`	Slope of the peak exercise ST	Categorical
`ca`	Number of major vessels (0-3)	Float
`thal`	Thalassemia (3 = Normal, 6 = F	Categorical

2.2 Target Variable Mapping

The raw num target in the UCI dataset represents angiographic disease status from 0 (no presence) to 4 (severe). In accordance with standard practice, this is binarized:

- 0 -- No presence of heart disease (original value 0)
- 1 -- Presence of heart disease (original values 1, 2, 3, or 4)

3. System Architecture

The following diagram illustrates the complete end-to-end architecture of the MLOps pipeline, from data acquisition to production monitoring:

[Architecture diagram - see report.md for Mermaid source]

3.1 Project Structure

```
Heart-Disease-Predictions-MLOPS/
|-- src/
|   |-- components/
|   |   |-- data_ingestion.py          # Fetches UCI dataset, splits train/test
|   |   |-- model_transformation.py    # Preprocessing pipeline (impute + scale + encode)
|   |   |-- model_trainer.py           # Trains models, logs to MLflow
|   |-- pipeline/
|   |   |-- train_pipeline.py          # Orchestrates full training workflow
|   |   |-- predict_pipeline.py        # Inference pipeline for serving
|   |-- config/
|   |   |-- configuration.py           # Centralized configuration
|   |-- utils.py                       # Serialization, evaluation helpers
|   |-- logger.py                      # Structured logging setup
|   |-- exception.py                   # Custom exception handling
|-- app.py                             # Flask REST API
|-- templates/
|   |-- index.html                     # Web UI for predictions
|-- static/
|   |-- style.css                      # UI styling
|-- notebooks/
|   |-- EDA.ipynb                     # Exploratory Data Analysis
|   |-- Model_Training.ipynb          # Model training notebook
|-- tests/
|   |-- test_data_ingestion.py         # Data ingestion tests
|   |-- test_prediction_schema.py      # Prediction schema tests
|-- artifacts/
|   |-- data/                         # Raw, train, test CSVs
|   |-- model.pkl                     # Best trained model
|   |-- preprocessor.pkl               # Fitted preprocessor
```

```
|   |-- plots/
|       |-- model_comparison.png      # Model comparison chart
|-- .github/workflows/
|   |-- ci.yml                        # GitHub Actions CI pipeline
|-- k8s/
|   |-- deployment.yaml               # Kubernetes Deployment
|   |-- service.yaml                  # Kubernetes Service
|-- monitoring/
|   |-- prometheus.yml                # Prometheus scrape configuration
|-- Dockerfile                       # Container build instructions
|-- requirements.txt                  # Python dependencies
|-- setup.py                          # Package setup
|-- environment.yml                   # Conda environment specification
```

4. Environment Setup & Installation

4.1 Prerequisites

- Python 3.11+
- Conda (Miniconda or Anaconda)
- Git
- Docker Desktop (for containerization)
- kubectl + Minikube or Docker Desktop Kubernetes (for K8s deployment)

4.2 Clone the Repository

```
git clone https://github.com/HARSH6355/Heart-Disease-Predictions-MLOPS.git
cd Heart-Disease-Predictions-MLOPS
```

4.3 Create and Activate the Conda Environment

Option A -- Using the existing local .conda environment (Windows PowerShell):

```
.\.conda\python.exe -m pip install -r requirements.txt
.\.conda\python.exe -m pip install -e .
```

Option B -- Creating a fresh Conda environment:

```
conda env create -f environment.yml
conda activate heart-disease-mlops
pip install -r requirements.txt
pip install -e .
```

4.4 Python Dependencies

The project relies on the following key libraries defined in requirements.txt:

Library	Purpose
`pandas`, `numpy`	Data manipulation and numerical
`scikit-learn`	Preprocessing, model training,
`matplotlib`, `seaborn`	Data visualization
`flask`	REST API web framework
`dill`	Object serialization (models a
`mlflow`	Experiment tracking and model
`ucimlrepo`	Programmatic dataset fetching
`pytest`	Unit testing framework
`prometheus-client`	Prometheus metrics instrumenta
`unicorn`	Production WSGI server
`ruff`	Python code linting

4.5 Train the Model

```
.\.conda\python.exe -m src.pipeline.train_pipeline
```

This generates the following artifacts:

- artifacts/data/raw.csv -- Full raw dataset
- artifacts/data/train.csv -- Training split (80%)
- artifacts/data/test.csv -- Test split (20%)
- artifacts/preprocessor.pkl -- Fitted preprocessing pipeline
- artifacts/model.pkl -- Best performing model
- artifacts/plots/model_comparison.png -- Visual model comparison
- mlruns/ -- MLflow experiment tracking data

4.6 Run the API

```
.\.conda\python.exe app.py
```

The API starts on <http://127.0.0.1:8080> with the following endpoints:

Endpoint	Method	Description
`/`	GET	Web-based prediction form
`/health`	GET	Health check (returns `{ "statu
`/predict`	POST	Accepts patient JSON, returns
`/metrics`	GET	Prometheus metrics endpoint

4.7 Launch MLflow UI

```
.\.conda\Scripts\mlflow.exe ui --backend-store-uri sqlite:///mlflow.db
```

Open the MLflow dashboard at <http://127.0.0.1:5000>.

5. Exploratory Data Analysis (EDA)

A comprehensive EDA was conducted in notebooks/EDA.ipynb to understand the dataset characteristics, identify patterns, and inform modelling decisions.

5.1 Data Loading & Metadata

The dataset was loaded programmatically from the UCI Machine Learning Repository using the `ucimlrepo` library. The metadata confirms 303 instances across 13 features with a single target column.

```
!Data loading and UCI metadata
```

5.2 Variable Information

Each feature's role (Feature/Target), data type (Integer, Categorical, Float), and missing value status was examined. Notable findings: `ca` has 4 missing values and `thal` has 2 missing values.

```
!Variable information and feature types
```

5.3 Data Shape & Schema

Using `df.info()`, the dataset shape of (303, 14) was confirmed. All 303 entries are present for most columns, with minor missing data in `ca` (299 non-null) and `thal` (301 non-null).

```
!DataFrame info showing shape and dtypes
```

5.4 Statistical Summary

`df.describe()` provided summary statistics across all numerical features. Key observations:

- Age ranges from 29 to 77 years, with a mean of 54.4 years
- Cholesterol ranges from 126 to 564 mg/dl
- Maximum Heart Rate ranges from 71 to 202 bpm
- Resting BP ranges from 94 to 200 mm Hg

The raw target distribution (`df.num.value_counts()`) showed: 164 (no disease), 55, 36, 35, and 13 across severity levels 0-4.

!Statistical summary and target distribution

5.5 Target Binarization

The multi-class target was binarized into a binary classification problem:

```
df['num'] = df['num'].apply(lambda x: 0 if x == 0 else 1)
df.rename(columns={'num': 'target'}, inplace=True)
```

After binarization, the dataset shape remained (303, 14) with balanced classes.

!Target binarization and column listing

5.6 Target Class Distribution

The class distribution is approximately balanced:

- No Heart Disease (0): 164 instances (54.1%)
- Heart Disease (1): 139 instances (45.9%)

This near-balanced distribution is advantageous as it eliminates the need for class imbalance techniques such as SMOTE or class weighting.

!Heart Disease Distribution bar chart

5.7 Correlation Analysis

A correlation heatmap was generated to identify relationships between features and the target variable:

!Correlation heatmap of all features

Key Correlation Findings:

Feature	Correlation with Target	Interpretation
`thal`	**+0.53**	Strongest positive association
`ca`	**+0.46**	Number of fluoroscopy-coloured
`exang`	**+0.42**	Exercise-induced angina is a s
`oldpeak`	**+0.42**	ST depression correlates with
`cp`	**+0.41**	Chest pain type is highly pred
`thalach`	**?0.42**	Higher max heart rate -> lower
`slope`	**+0.33**	Peak exercise ST segment slope

Additionally, no severe multicollinearity was observed among predictors, with the highest inter-feature correlation being 0.58 between oldpeak and slope.

!Correlation analysis insights and age distribution code

5.8 Feature vs. Target Analysis -- Chest Pain Type

Analysis of chest pain types revealed a strong association with heart disease status:

- Asymptomatic (cp=4): Highest prevalence of heart disease
- Typical Angina (cp=1), Atypical Angina (cp=2), Non-Anginal Pain (cp=3): More frequently observed among individuals without heart disease

These findings confirm that chest pain type is an important predictive feature for heart disease diagnosis.

!Chest Pain Type vs Heart Disease -- code and chart

!Chest Pain Type vs Heart Disease -- full bar chart with interpretation

6. Feature Engineering & Preprocessing

Feature transformation is implemented in `src/components/model_transformation.py` using scikit-learn's `ColumnTransformer` to ensure a reproducible, leakage-free preprocessing pipeline.

6.1 Feature Categorization

Features are split into two groups:

Numerical Features: age, trestbps, chol, thalach, oldpeak

Categorical Features: sex, cp, fbs, restecg, exang, slope, ca, thal

6.2 Preprocessing Pipeline

```
# Numerical pipeline
numeric_pipeline = Pipeline([
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler()),
])

# Categorical pipeline
categorical_pipeline = Pipeline([
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("one_hot_encoder", OneHotEncoder(handle_unknown="ignore")),
])

# Combined preprocessor
preprocessor = ColumnTransformer([
    ("num_pipeline", numeric_pipeline, numerical_columns),
    ("cat_pipeline", categorical_pipeline, categorical_columns),
])
```


6.3 Design Decisions

Decision	Rationale
Median imputation for nume	Robust to outliers in clinical
Mode imputation for catego	Preserves the most common cate
StandardScaler for numeric	Z-score normalization aids con
OneHotEncoder with `handle	Handles unseen categories duri
Artifact serialization (`p	Ensures exact same transformat

6.4 Data Splitting Strategy

- Train/Test Ratio: 80/20 split
- Stratification: Stratified on the target column to maintain class proportions in both splits
- Random State: Fixed at 42 for reproducibility

7. Modelling Choices & Evaluation

Two classification models were trained and compared in `src/components/model_trainer.py`. The evaluation strategy combines 5-fold Stratified Cross-Validation on the training set with a final evaluation on the held-out test set.

7.1 Model Configurations

Logistic Regression:

- `max_iter=1000` -- Ensures convergence with high-dimensional one-hot encoded features
- `random_state=42` -- Reproducibility

Random Forest Classifier:

- `n_estimators=200` -- Sufficient ensemble size for a 303-sample dataset
- `max_depth=8` -- Prevents overfitting on a small dataset
- `min_samples_leaf=2` -- Additional regularization
- `random_state=42` -- Reproducibility

7.2 Evaluation Metrics

The following metrics were computed for comprehensive model assessment:

Metric	Why It Matters
Accuracy	Overall correctness
Precision	Proportion of positive predict

Recall	Proportion of actual positives
F1-Score	Harmonic mean of precision and
ROC-AUC	Discrimination ability across

7.3 Results Comparison

Model	Split	Accuracy	Precision	Recall	F1-Score	ROC-AUC
Logistic Regression	5-Fold CV (Mean)	0.843	0.869	0.774	0.817	0.894
	Hold-out Test	**0.885**	**0.839**	**0.929**	**0.881**	**0.966**
Random Forest	5-Fold CV (Mean)	0.834	0.835	0.800	0.815	0.894
	Hold-out Test	0.852	0.806	0.893	0.847	0.922

7.4 Model Comparison Visualization

!Heart Disease Model Comparison -- ROC-AUC and Recall

7.5 Model Selection Rationale

Logistic Regression was selected as the production model based on:

1. Superior ROC-AUC (0.966 vs. 0.922): Demonstrates better discrimination ability across all probability thresholds.
2. Higher Recall (0.929 vs. 0.893): In a healthcare context, minimizing false negatives (missed heart disease cases) is critical. Logistic Regression correctly identifies 92.9% of actual heart disease cases.
3. Better F1-Score (0.881 vs. 0.847): Best balance of precision and recall.
4. Interpretability: Logistic Regression provides interpretable coefficients, enabling clinicians to understand which features drive predictions.
5. Simplicity: Fewer hyperparameters reduce the risk of overfitting on this relatively small dataset.

The selected model is serialized to artifacts/model.pkl using dill for production deployment.

8. Experiment Tracking with MLflow

All experiments are tracked using MLflow with a SQLite backend (mlflow.db). Every training run logs:

- Parameters: Model name, hyperparameters (e.g., max_iter, n_estimators, max_depth)
- Metrics: Accuracy, Precision, Recall, F1-Score, ROC-AUC (both CV and test set)
- Artifacts: Trained model binaries, comparison plots

8.1 MLflow Experiment Dashboard

The heart-disease-classification experiment contains three runs:

1. Logistic Regression -- Full training with all metrics
2. Random Forest -- Full training with all metrics
3. best-model-summary -- Summary run logging the winner and the comparison chart

!MLflow experiment runs showing all three experiments

8.2 Metrics Comparison View

MLflow's built-in chart view enables side-by-side comparison of all metrics across runs:

!MLflow metrics comparison -- accuracy, cv_accuracy, cv_f1, cv_precision, cv_recall, cv_roc_auc

8.3 How to Reproduce

```
# Run the training pipeline (logs are created automatically)
.\.conda\python.exe -m src.pipeline.train_pipeline

# Launch MLflow UI to inspect experiments
.\.conda\Scripts\mlflow.exe ui --backend-store-uri sqlite:///mlflow.db
```

Navigate to <http://127.0.0.1:5000> -> Select the heart-disease-classification experiment -> Compare runs.

9. Flask Prediction API & Web Interface

The serving layer is a Python Flask application (app.py) that loads the trained model and preprocessor to serve real-time predictions via REST endpoints and a web-based form.

9.1 API Endpoints

Endpoint	Method	Description	Response
`/`	GET	Serves the prediction web form	HTML page
`/health`	GET	Health check	`{"status": "ok"}`
`/predict`	POST	Accepts patient data, returns	`{"prediction": 0/1, "confiden
`/metrics`	GET	Prometheus metrics	Plain text metrics

9.2 Prediction Flow

1. User submits patient data via the web form or API call
2. Flask receives the JSON payload and validates all 13 required fields

3. Input is converted to a HeartDiseaseInput dataclass and then to a DataFrame
4. The preprocessor (preprocessor.pkl) transforms features identically to training
5. The model (model.pkl) generates a prediction and confidence score
6. Response is returned with prediction class, confidence, and human-readable risk label

9.3 Web Interface

The web UI provides a modern, dark-themed form for entering patient data:

!Heart Disease Risk Prediction -- Web Form

9.4 Prediction Result

After submitting patient data, the result card shows the prediction with a confidence bar:

!Prediction result showing Heart Disease Risk Detected with 56.3% confidence

9.5 Sample API Request

```
Invoke-RestMethod -Uri http://127.0.0.1:8080/predict `
-Method Post `
-ContentType "application/json" `
-Body '{"age":54,"sex":1,"cp":4,"trestbps":130,"chol":250,"fbs":0,"restecg":1,"thalach":150,"exang":0,"oldpeak":1.0}'
```

Sample Response:

```
{
  "prediction": 1,
  "confidence": 0.563,
  "risk_label": "heart disease risk"
}
```

10. Unit Testing

Automated tests are implemented using Pytest to validate critical components.

10.1 Test Coverage

Test File	What It Tests
`tests/test_data_ingestion.py`	Validates that the data ingest
`tests/test_prediction_schema.	Validates the `HeartDiseaseInp

10.2 Running Tests

```
.\.conda\python.exe -m pytest -q tests
```

All tests pass successfully, confirming the integrity of the data ingestion and prediction schema components.

11. CI/CD Pipeline -- GitHub Actions

Continuous Integration is implemented via GitHub Actions in `.github/workflows/ci.yml`. The pipeline automates testing, linting, training, and artifact collection on every push or pull request.

11.1 Workflow Configuration

```
name: mlops-ci

on:
  push:
    branches: ["main", "master"]
  pull_request:
    branches: ["main", "master"]
  workflow_dispatch:

jobs:
  test-and-train:
    runs-on: ubuntu-latest
    steps:
      - Checkout repository
      - Set up Python 3.11
      - Install dependencies (pip install -r requirements.txt && pip install -e .)
      - Run unit tests (pytest -q)
      - Run linting (ruff check .)
      - Run training pipeline (python -m src.pipeline.train_pipeline)
      - Upload generated artifacts (artifacts/, logs/, mlruns/)
```

11.2 GitHub Actions -- Workflow Runs

The CI pipeline runs successfully on both push events and manual triggers:

!GitHub Actions showing 2 successful mlops-ci workflow runs

11.3 GitHub Actions -- Job Steps

All steps in the test-and-train job complete successfully:

Step	Duration	Status
Set up job	0s	?
Checkout repository	2s	?
Set up Python	0s	?
Install dependencies	1m 15s	?
Run unit tests	2s	?
Run linting	0s	?
Run training pipeline	22s	?
Upload generated artifacts	2s	?

!GitHub Actions job detail showing all steps passing

12. Docker Containerization

The application is containerized using a production-ready Dockerfile for consistent deployment across environments.

12.1 Dockerfile

```
FROM python:3.11-slim

WORKDIR /app

ENV PYTHONDONTWRITEBYTECODE=1
ENV PYTHONUNBUFFERED=1

COPY requirements.txt setup.py ./
COPY src ./src
COPY templates ./templates
COPY static ./static
COPY artifacts ./artifacts
COPY app.py ./

RUN pip install --no-cache-dir --upgrade pip \
    && pip install --no-cache-dir -r requirements.txt \
    && pip install --no-cache-dir -e .

EXPOSE 5000

CMD ["gunicorn", "--bind", "0.0.0.0:5000", "app:app"]
```

12.2 Build & Run Commands

```
# Build the Docker image
```

```
docker build -t heart-disease-api:latest .

# Run the container
docker run -p 5000:5000 heart-disease-api:latest
```

12.3 Docker Desktop -- Image Built

The heart-disease-api:latest image was successfully built and is visible in Docker Desktop:

!Docker Desktop showing heart-disease-api image (3.32 GB)

12.4 Docker Container -- Running

The container runs with Gunicorn (production WSGI server) and logs confirm successful startup:

!Docker Desktop container logs showing Gunicorn startup

12.5 Key Design Decisions

Decision	Rationale
`python:3.11-slim` base image	Minimal footprint while mainta
`PYTHONDONTWRITEBYTECODE=1`	Prevents `.pyc` file generatio
`PYTHONUNBUFFERED=1`	Ensures logs are flushed immed
`gunicorn` instead of Flask de	Production-grade, concurrent r
Pre-trained artifacts copied i	No training required at runtim

13. Kubernetes Deployment

For production-grade orchestration, Kubernetes manifests are provided in the k8s/ directory.

13.1 Deployment Manifest (`k8s/deployment.yaml`)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: heart-disease-api
spec:
  replicas: 1
  selector:
    matchLabels:
      app: heart-disease-api
  template:
```

```
spec:
  containers:
    - name: heart-disease-api
      image: heart-disease-api:latest
      ports:
        - containerPort: 5000
      readinessProbe:
        httpGet:
          path: /health
          port: 5000
        initialDelaySeconds: 10
        periodSeconds: 10
      livenessProbe:
        httpGet:
          path: /health
          port: 5000
        initialDelaySeconds: 20
        periodSeconds: 20
```

13.2 Service Manifest (`k8s/service.yaml`)

```
apiVersion: v1
kind: Service
metadata:
  name: heart-disease-api-service
spec:
  type: LoadBalancer
  selector:
    app: heart-disease-api
  ports:
    - protocol: TCP
      port: 80
      targetPort: 5000
```

13.3 Deploy to Kubernetes

```
kubectl apply -f k8s/
kubectl get pods
kubectl get svc
```

13.4 Kubernetes -- Deployment Results

Successful deployment showing the pod running and service exposed via LoadBalancer:

!kubectl apply, get pods, and get svc showing running deployment

Deployment Status:

- Pod: heart-disease-api-cb4c58f7c-rpf9z -- 1/1 Running, 0 restarts

- Service: heart-disease-api-service -- LoadBalancer, External IP: localhost:30364

13.5 Application Running on Kubernetes

The application served through Kubernetes on localhost:

!Application running on localhost via Kubernetes LoadBalancer

13.6 Prediction via Kubernetes

Making a prediction through the Kubernetes-served application:

!Heart Disease Risk Detected result via Kubernetes deployment

13.7 Health Probes

Probe Type	Path	Initial Delay	Period
Readiness	/health`	10s	10s
Liveness	/health`	20s	20s

These probes ensure Kubernetes automatically restarts unhealthy pods and only routes traffic to ready instances.

14. Monitoring with Prometheus

The API is instrumented with Prometheus metrics using the prometheus-client Python library for production observability.

14.1 Instrumented Metrics

Metric Name	Type	Description	Labels
heart_api_requests_total`	Counter	Total API requests	endpoint, method, http_status
heart_api_request_latency_sec	Histogram	Request latency in seconds	endpoint
heart_api_predictions_total`	Counter	Total predictions by class	prediction (0 or 1)

14.2 Prometheus Configuration

The monitoring/prometheus.yml configuration scrapes the Flask app's /metrics endpoint:

```
global:
  scrape_interval: 15s
```

```
scrape_configs:
  - job_name: "heart-disease-api"
    metrics_path: /metrics
    static_configs:
      - targets: ["host.docker.internal:5000"]
```

14.3 Accessing Metrics

The metrics endpoint at /metrics exposes Prometheus-formatted data:

```
# HELP heart_api_requests_total Total API requests
# TYPE heart_api_requests_total counter
heart_api_requests_total{endpoint="/predict",method="POST",http_status="200"} 5.0

# HELP heart_api_request_latency_seconds API request latency in seconds
# TYPE heart_api_request_latency_seconds histogram
heart_api_request_latency_seconds_bucket{endpoint="/predict",le="0.005"} 0.0
heart_api_request_latency_seconds_bucket{endpoint="/predict",le="0.01"} 3.0
...

# HELP heart_api_predictions_total Total predictions by class
# TYPE heart_api_predictions_total counter
heart_api_predictions_total{prediction="1"} 3.0
heart_api_predictions_total{prediction="0"} 2.0
```

14.4 Monitoring Architecture

[Architecture diagram - see report.md for Mermaid source]

This setup enables:

- Request rate monitoring -- Track API traffic patterns
- Latency percentiles -- Identify performance degradation
- Prediction drift detection -- Monitor class distribution changes over time

15. Conclusion & Future Work

15.1 Summary

This project demonstrates a complete, production-ready MLOps lifecycle for heart disease risk prediction:

Component	Implementation
Data Acquisition	Automated from UCI Repository

EDA	Comprehensive analysis with 11
Preprocessing	Reproducible scikit-learn `Col
Model Training	Logistic Regression vs. Random
Experiment Tracking	MLflow with SQLite backend
Model Serving	Flask REST API with web interf
Testing	Pytest unit tests
CI/CD	GitHub Actions (test -> lint ->
Containerization	Docker with Gunicorn
Orchestration	Kubernetes with health probes
Monitoring	Prometheus metrics instrumenta

The best performing model (Logistic Regression) achieved an F1-Score of 0.881 and ROC-AUC of 0.966 on the test set.

15.2 Future Improvements

1. Model Enhancement: Explore gradient boosting methods (XGBoost, CatBoost) and hyperparameter optimization with Optuna
2. Data Drift Monitoring: Integrate Evidently AI for automated data and model drift detection
3. Model Registry: Leverage MLflow Model Registry for versioned model deployment with staging/production stages
4. Grafana Dashboards: Create visual dashboards for real-time monitoring of prediction latency, request rates, and class distributions
5. A/B Testing: Implement canary deployments to compare model versions in production
6. Authentication: Add API key or OAuth2 authentication for the prediction endpoint
7. Batch Prediction: Support batch inference for large-scale patient screening

Code Repository

GitHub: <https://github.com/HARSH6355/Heart-Disease-Predictions-MLOPS>

Report generated for the Heart Disease Prediction MLOps Assignment -- July 2026